

Portage - A look under the surface

dev.gentoo.org/~genone/docs/fosdem-2007-talk.pdf

Marius Mauch <genone@gentoo.org>

Free and Open Source Developers European Meeting,
2007



Contents

- Recently added features
- The EAPI problem
- Backend technologies
- Popular requests and their problems
- Discussion



Recently added features

elog

Resolver improvements

Manifest2

License filtering

Per-package default USE flags

Others



Where to find new features

- `/usr/share/doc/portage-$version/NEWS.gz`
- sticky topics on the “Portage and Programming” forum
- GWN
- gentoo-portage-dev mailing list
- gentoo-dev mailing list
- #gentoo-portage IRC channel



elog: What is elog?

- A framework for logging ebuild generated messages
- An ebuild function to send messages to the user
- A set of modules to dispatch messages to the user



elog: What is elog?

- A framework for logging ebuild generated messages
- An ebuild function to send messages to the user
- A set of modules to dispatch messages to the user
- A solution to an annoying problem



elog: How does it work?

1. Ebuild sends a message via einfo/elog/ewarn/error
2. At merge time of each package portage collects these messages
3. Collected messages are filtered and relayed to dispatch modules
4. Dispatch modules process the messages
5. At exit time portage calls finalization code of dispatch modules



elog: Available modules

save Saves messages in files (one file per package)

mail Sends messages by mail (one mail per package)

syslog Inject messages into the syslog facility

custom Pass messages to a custom logging tool/script

??? New modules are easy to create



elog: Problems

- Addition of new preferred loglevel
- Not compatible with old enotice hack
- No "official" tools for handling logfiles
- Sometimes confusing configuration
- Messages not always in chronological order



Resolver improvements: Blockers

- Blockers from installed packages are now also considered
 - `DEPEND="!foo"` in package `bar` automatically adds `DEPEND="!bar"` in package `foo`
- In some cases blockers can be handled automatically without manual unmerge
 - Only works when the blocker is made irrelevant by an update in the same graph



Resolver improvements: Full dependency graph

- Enables detection of circular dependencies
- Avoids up-/downgrade cycles
- Ensures correct merge order
- Includes multiple SLOTS of a package



Manifest2: Overview

- A new container format for ebuild related checksums
- Expected to replace the old digest system in the near future
- Fully backwards compatible
- Adds more control about checksum verification



Manifest2: Motivation

The old checksum system was very inefficient:

- Filename and size had to be repeated for each checksum
- Needed many tiny files - high filesystem overhead
- All files equal

The new system handles this much better:

- Each file is only listed once
- All checksums are stored in a single file (per package)
- Allows specifying a filetype that can be used for filtering



License filtering: What is that?

- More commonly known as ACCEPT_LICENSE
- New visibility filter on top of keyword filtering (later)
- Replacement for check_license ebuild function



License filtering: What is that?

- More commonly known as ACCEPT_LICENSE
- New visibility filter on top of keyword filtering (later)
- Replacement for check_license ebuild function
- Implemented but not yet merged



License filtering: Motivation

- People want to build “free” systems
- Replace `check_license` with a better system
- Generally make use of the `LICENSE` variable



Default IUSE flags

- Allows recommending USE flags per package
- Extends IUSE syntax with **Example**: `+flag`
- Needs EAPI bump before usage (later)
- Lowest config layer - overridden by all other layers



Per profile package.use

- Like normal package.use
- Useful for fine-graining profile use flags
- Overridden by flags in make.conf (unlike normal package.use)
- Less compability issues than default IUSE flags



Directories-as-files: WTF does that mean?

- Most config files can be splitted in multiple files
- Allows a more flexible organization
- Files can also be symlinks - useful for shared configuration
- Can reserve filenames for special uses
- Also supports subdirectories



SLOT dependencies

- Allows to match only ebuilds in a specific SLOT
- Syntax: **Example**: `foo/bar:2` to only select ebuilds with SLOT=2
- Possible replacement for version ranges in some cases
- Needs EAPI bump before usage in ebuilds (later)



New FEATURES

parallel-fetch performs downloading and compiling in parallel

userfetch downloads distfiles as unprivileged user

stricter Makes QA checks more severe (aborts build if checks fail), not intended for normal usage

installsources Install sourcecode as well as compiled files

splitdebug Store debug symbols in separate files when stripping binaries (needs installsources)



Others

per package use.mask Allows profiles to mask USE flags
per package

forced USE flags Allows profiles to force USE flags as
“on”

extended versioning syntax Support for cvs versions and
multiple suffixes

bashrc hooks Can run custom code before/after each
ebuild phase

rsync option handling New variables to send options to
rsync in a generic way



The EAPI problem

EAPI description

EAPI problems



EAPI: Motivation

- Introduction of new features
 - Changes to ebuild phases (src_configure)
 - Metadata syntax changes (default IUSE)
 - New ebuild functions (elog)
- Removing legacy stuff
 - RESTRICT=noXXX
 - PROVIDE entries



EAPI: Implementation

- New metadata key
- Packages with unsupported EAPI value are masked
- Switching semantics based on EAPI value



EAPI: Implementation

- New metadata key
- Packages with unsupported EAPI value are masked
- Switching semantics based on EAPI value
- Mostly untested so far



EAPI: Limitations

- Limited to ebuild internals
- Not only ebuilds need versioning (repo layout, support files)
- No clear syntax definition
- Alternative: Repository versioning



EAPI: Problems

- Current semantics are mostly undefined (EAPI=0)
- Planned semantics are undefined (EAPI=1)
- Exact syntax is undefined
- Current implementation is limited to visibility filtering



Backend technologies

The visibility system

The configuration stack

The cache system



Visibility system: Concept

- Usually referred to as “masking”
- Core: All ebuilds available
- Core selection is processed through boolean filters
- Multiple layers stacked on each other
- Visibility vs. Grouping



Visibility system: Layers

Corruption if metadata is unreadable

EAPI if ebuild version isn't supported

Profile if package is incompatible with profile

Package.mask general purpose (configurable)

Keywords Stability/Age (configurable)

License if license is unacceptable (configurable)

- potentially others



Configuration stack: Concept

- Portage configuration assembled from multiple locations
- Files and variables
 - Incremental and Override variables
 - Files are generally incremental
- Locations define responsibility
- Non-unified semantics



Configuration stack: Variable Layers

Make.globals Portage defaults (Portage developers)

Make.defaults Profile settings (profile maintainers,
Releng)

Make.conf Global user configuration (user)

Environment General override (user, tools)



Configuration stack: File Layers

`$PORTDIR/profiles` Global repository settings (ebuild
devs)

`/etc/make.profile` Profile settings (profile maintainers,
Releng)

`/etc/portage/profile` Profile overrides (user)

`/etc/portage` User settings (user)

Not all files available in all locations!



Cache system: Concept

- Motivation: ebuild parsing is slow
- Parsed metadata variables/values are cached
- Modular



Cache system: Restrictions

- Values must be constant
 - Metadata variables can't be affected by local configuration
- Cache must be checked for staleness
 - Complex due to eclasses
 - Makes access rather slow
- Access patterns are fixed → less optimization potential



Popular requests and their problems

Database backend

Reverse dependency support

USE dependencies

Remote ebuild tree

GPG support



Why not use a database backend?

- Usual problem: way to vague
 - Backend for what?
 - What exactly is a “database”?
- General misassumption: database == fast



Problems of a RDBMS cache backend

- Additional dependencies
- No practical benefit
 - data access patterns aren't designed with a RDBMS in mind
 - portage already does it's own caching
- sqlite/mysql modules already exist(ed)



Portage doesn't have reverse dependency support

- “Reverse dependencies” aren’t a feature but a property of a graph edge
- People refer to different features here:
 - Automatic rebuilding of link-level dependencies (`revdep-rebuild`)
 - Recursive unmerging of dependencies (`emerge -depclean`)
 - Unmerging of reverse dependencies
 - Prevent unmerging of dependencies



Portage lacks USE dependencies

- Most wanted feature by ebuild developers
- Just nobody has implemented it yet in portage
- People are scared by resolver code apparently



Why can't portage download the ebuilds when it needs them?

- Lack of atomicity
- Possible integrity issues
- Requires to be always online
- Lack of container format
- Harder to support
- Metadata has to be fetched completely anyway



What's the status of GPG support in portage?

- Original plans date back to 2003 and earlier
- No real specification was ever made
- No key policy was ever created
- Some prototype code was added to portage in 2004
- Current GPG support is very incomplete and unmaintained
- A new specification is in the works (hopefully)



Discussion

